

Chapter 1

Case study: Zipf's Law and Music

Max Ward (20748588)

1.1 Music as Language

Zipf's law was introduced in Chapter 5. To recapitulate, it posits that the frequency of a word is inversely proportional to its rank. The rank of a word being defined as its index in a frequency table. The canonical example of this is the words of long piece of literature, usually a novel. It is a classical example of a scale-free distribution, and is defined by $f = r^{-a}$ where f is frequency and r is rank. A frequency distribution is said to be Zipfian if $a \approx 1$. Because of this, when plotted on a log-log scale (equivalent to taking the log of both sides), a Zipfian distribution appears to be a line with slope $m \approx -1$. Though Zipf's law is typically described in terms of novels, and thus written language, it actually refers to any natural language. Natural language is generally regarded to be any non-contrived language—which is to say that its use and invention is unpremeditated. Clearly most written and spoken languages are natural languages, however, written and performed music is also a type of natural language.

Figure of a log log zipf graph here

Music, among other art forms, is often regarded as communication. Like other forms of human communication it has structure and rules. While I shall not give a formal definition of these, the reader may refer to http://en.wikipedia.org/wiki/Music_theory for a concise overview. This chapter investigates Zipf's law in music. This is done by using the MIDI file type, which is a relatively simple and widely used way to encode digital music.

Exercise 1.1. *There are specific parts of the human brain involved in language. For example, damage to Broca's area can profoundly affect ones ability to understand spoken and written language. Does an analogous region exist for processing music? (Hint: try searching for 'Amusia'.)*

1.2 The MIDI Format

The MIDI file format comprises a header and a series of one or more tracks. Every MIDI track also contains a header, which is followed by a sequence of events. There are many event types which refer to any time related possibility. However, the only two which we consider in this investigation are the 'Note On' and 'Note Off' events. A Note On event has a time, a velocity, and a note. The time is when it is played, the velocity is the force with which it is played, and the note is the specific sound frequency of the event. It should be noted that velocity is not volume. It is possible that a MIDI instrument could have a linear velocity-volume relationship, but it could be exponential, logarithmic, or anything else. There are 128 distinct notes in the MIDI format, these are analogous to an extended piano (which has only 88 keys), with each increment indicating a semitone increase in sound frequency.

Figure of a piano keyboard here.

Parsing the MIDI file format can be difficult, as such, it is often relegated to a library. Python does not come with an easy to install MIDI parsing library. However, an open source effort called `python-midi` is available on Github under the MIT licence. It can be found here <https://github.com/vishnubob/python-midi>. Download or clone the repository and use the source files in your project. Once they are in the same folder, you should be able to import them like any Python library. I have provided a basic example that loads a MIDI file, given a file name from the command line. The `fileio` module is part of `python-midi`. It contains functions related to reading and parsing MIDI files.

```
import sys
import fileio

def main(name, fname, *args):
    pattern = fileio.read_midifile(fname)
    print pattern

if __name__ == '__main__':
    main(*sys.argv)
```

The way MIDI patterns are represented in `python-midi` is straightforward. A pattern is a list of tracks. A track is a list of events. Every event has a type, and associated metadata. All of these classes also have the `__repr__` method overloaded. This makes it easy to determine what information is available. A quirk of some older MIDI files is that a Note On event with velocity 0 is equivalent to the corresponding Note Off event. Keep this in mind when experimenting with MIDI files.

The Classical Archive has some excellent MIDI encodings of classical music (available at <http://www.classicalarchives.com/midi.html>). You may also wish to find some songs which you enjoy, many people have made their MIDI versions of songs freely available on the Internet. There is also software that allows the user to write MIDI music using a graphic environment. If you are so inclined, it is possible to analyse your own compositions.

Exercise 1.2. *Research the MIDI format. How hard do you think it would be to write your own parser? Remember, if you improve `python-midi`, submit a pull request to Github!*

Exercise 1.3. *Attempt to write a script that outputs the number of distinct notes in a MIDI file. My solution is available at [URL](#).*

1.3 Words and Music

Zipf's law refers to the frequency of words found in a large corpus of natural language. It is hard to see how the concept of a word could apply to musical compositions. Of course, any interpretation of a word in music is necessarily arbitrary and somewhat open to interpretation. Despite this, I shall now describe several feasible candidates. Clearly the most obviously irreducible component of a musical composition is the note. Building upon this starting point, one might analyse the velocities, lengths, or the spacing between similar notes. More complex interpretations may attempt to group notes into structures like chords or melodies. I shall, however, start with a much simpler interpretation.

1.3.1 Notes

As I have said, we can break a piece of music up into notes. If we define a word as a single note, we can then collect the frequencies of these notes, and thence subject them to Zipfian analysis. This necessitates going through every MIDI event in every MIDI track for any given pattern. The `python-midi` package makes this quite easy, as tracks and events are represented as objects that inherit from `list`. Here is a simple function that returns all the notes (in no particular order) given the path to a MIDI file.

```
def get_notes(fname):
    pattern = fileio.read_midifile(fname)
    notes = []
    for track in pattern:
        for event in track:
            if event.name == "Note On" and event.data[1] > 0:
                notes.append(event.data[0])
    return notes
```

The first line reads and parses the MIDI into a pattern. The follow two loops examine every event in the pattern. Every time a genuine Note On event is spotted (remember that a Note On event with velocity 0 is actually a Note Off event) it is added to the `notes` list. It should be pointed out that event data is stored in the `event.data` member. For Note On events, index zero is the MIDI note, and index one is the velocity. After executing this function, we now have everything we need to calculate frequency as a function of rank.

1.3.2 Consecutive Notes

A more versatile way of using note frequencies is to consider sequences of consecutive notes. This is equivalent to sorting all the notes in the MIDI file by order of appearance,

then running a sliding window of fixed size over all the notes. This will result in a much larger corpus of possible words. For example, a window size of two will have 128^2 possible words, every combination of two MIDI notes. Having a window size of one is the same as the approach previously described in Section 1.3.1. Having a too large window size can be a problem, however, as it often results in having many words with very low frequency. I found that a window size of two was typically the best. Very long pieces of music may work with larger window sizes, however.

Figure of sliding window

Exercise 1.4. *See if you implement the sliding window note analysis technique described in this section. Note On objects have a member called `tick`, which is the number of ticks before they appear in the MIDI song. You may wish to use this for sorting. My solution can be found at [URL](#).*

1.3.3 Note and Velocity Pairs

I have outlined how we could use the occurrence of notes as an interpretation of words for Zipfian analysis. However, a Note On event has two components; the note and its velocity. Another reasonable interpretation of a word is to use a tuple of these data. A quirk of the `python-midi` package is that `event.data` is actually a list. Because of this, I have used the built-in function `tuple` to convert it to a tuple in the code snippet presented below.

```
def get_notevelocities(fname):
    pattern = fileio.read_midifile(fname)
    notes = []
    for track in pattern:
        for event in track:
            if event.name == "Note On" and event.data[1] > 0:
                notes.append(tuple(event.data))
    return notes
```

Exercise 1.5. *Can you think of any problems with this method? Would using a sliding window over note and velocity pairs be a good idea?*

1.3.4 Note Intervals

Thus far, all the word interpretations I have described deal with discrete events directly. As an alternative, we may wish to take a more high level approach. For every MIDI note frequency that appears in a MIDI pattern, the intervals between specific notes could be recorded. For example, if the note code 60 appears at tick 110, then later at 120, then again at 135, we would record intervals of 10 and 15. Such intervals could be collected for every note in the MIDI pattern, and collated into one list. We can then calculate the rank and frequency of note intervals for further analysis.

Exercise 1.6. *Implement the `get_noteintervals` function. Modifying the `get_notes` function might make this easier. Don't forget to sort all MIDI events by tick first. (Hint: use a dictionary to store the last occurrence of every note code.)*

Exercise 1.7. *Can you think of any other aspects of music that might reasonably conform to Zipf's law? If you can think of any, implement a getter function like those present in this section.*

1.4 Scale-free Music

In Chapter 5, we created a number of useful classes and functions for dealing with Zipf, and other power law, distributions. A copy of these files was helpful in analysing the words described in the preceding section. Using the `Pmf.py` and `Zipf.py` modules, a histogram can be created for ranks and frequencies. From here we can plot a log-log graph of ranks versus frequencies. If the graph looks somewhat like a straight line with slope ≈ -1 then this confirms that the rank-frequency function is Zipfian. For analysis I chose ‘Wish You Were Here’ and ‘Comfortably Numb’ by Pink Floyd, Rachmaninov’s ‘Allegro Ma Non Tanto’ (No. 3 in D minor, Op. 30), and Brahms ‘Hungarian Dance’ No. 1. First, let us see a simple distribution of single note frequencies.




Figure 1.1: Hungarian Dance No. 1 using single notes




Figure 1.2: Wish You Were Here using single notes



Figure 1.3: Allegro Ma Non Tanto using single notes




Figure 1.4: Comfortably Numb using single notes

Figures 1.1 - 1.4 show log-log plots of frequency versus rank for my selected songs. These graphs only look somewhat Zipfian. Typically a graph starts like a line, but drops off quickly towards the end. It is important to remember that songs are generally only a few minutes long, and are not comparable in size to a novel. As such, I argue that this constitutes evidence that music exhibits Zipf’s law to a remarkable degree. To further test my hypothesis, I ran the exact same test, except using a sliding window of size two. This means that every pair of adjacent notes was used a ‘word’ for the purposes of calculating a frequency.




Figure 1.5: Hungarian Dance No. 1 using a sliding window over pairs of notes




Figure 1.6: Wish You Were Here using a sliding window over pairs of notes




Figure 1.7: Allegro Ma Non Tanto using a sliding window over pairs of notes




Figure 1.8: Comfortably Numb using a sliding window over pairs of notes

As can be seen in Figures 1.5 - 1.8, using a sliding window over note pairs was a little more compelling. A similar pattern emerges: a line that eventually drops off towards the end. However the suddenness of the drop off is typically less, and the extend of the line typically larger. Brahms Hungarian Dance No. 1 in particular looks notably like a typical Zipf distribution.

Exercise 1.8. Use some music of your choosing to recreate these graphs. Do you see a similar pattern?

Exercise 1.9. In the previous chapter, several other interpretations of a musical 'word' were discussed. Test if they follow Zipf's law in the MIDI tracks you chose.

1.5 The Beauty of Context

In 1955, Simon published a paper which attempted to explain Zipf's law. He noted that it had been observed in many vastly different, and probably unconnected, domains. For example, Zipf's law can be found the distribution of scientists and their published papers, and as we discovered in Chapter 5, in the rank of cities by population. A copy of Simon's original paper, "On a Class of Skew Distribution Functions", can be found here <http://>

www.uvm.edu/~pdodds/teaching/courses/2009-08UVM-300/docs/others/1955/simon1955.pdf. It contains a lengthy and somewhat impenetrable explanation of why Zipf's law might exist. The essence of his idea is quite straightforward, however.

Simon posits that large bodies of literature of generally constructive. Consecutive words are appended one by one to the end, which is precisely what is happening as I write this. As words are added, a context is created. As this context emerges ex nihilo, it favours the reuse of previous words to maintain the context. In a similar sense, the introduction of new words is also inhibited. This model presupposes that context is an essential part of communication, and any other process that exhibits a scale-free distribution. Intuitively it is easy to see the importance of context to human communication. Say that we take a novel and randomly shuffle sentences around. It is overwhelmingly likely that the book is now unintelligible. The linear construction of the book's context would have been destroyed. Similarly, take a piece of music and shuffle around movements randomly. Though each individual component would make sense, the cohesion of the whole would be destroyed.

As recently as 2005, Bill, Juan, and Penousal used Zipf's law to define a fitness function for judging music. Obviously such a function is useful for the computer generation of pleasant music. They used a large array of tests, encompassing and eclipsing those presented here. Much like we found in the previous chapter, they discovered that Zipf's law was present in musical compositions. Furthermore, they found that it was not present in non-music compositions. In addition to this, they were able to train an artificial neural network to differentiate between genres of music, using the slope of the regression line for the log-log diagram of frequency and rank. It was also able to accurately distinguish the work of different composers within the same genre. On the strength of their findings, they concluded that Zipf's law was a necessary, though not complete, precondition for aesthetically pleasing music. The interested reader can find a copy of their paper, "Zipf's Law, Music Classification, and Aesthetics", at <http://www.idi.ntnu.no/emner/tdt04/2010/%20articles/Zipf-music-johannes.pdf>.

Clearly, claiming that music and language are similar because they both exhibit Zipf's law is an argument from analogy, and thus is also an holistic argument. However, what about Simon's model? Read his paper, does it appear to be reductionist, or holistic? Furthermore, is Simon's constructive argument based on an instrumentalist point of view, or realist? To clarify, we have two systems, that of language, and that of music. Both have the same property, Zipf's law. A holist would argue that this is because such a law is a fundamental or essential property of both. However, Simon has argued that this property can be explained constructively; at a very low level, music and language are similar, and as a result they share an emergent property. Sometimes the lines between holism, reductionism, instrumentalism, and realism can be ill-defined, this is a good example. Regardless, it is clear that Zipf's law is present in music. Why is another, and perhaps more intractable, question.

Exercise 1.10. Find examples of white noise, pink noise, and brown noise on the Internet. (YouTube is a good place to look.) Which do you prefer, and why? (Hint: Pink noise is also called $1/f$ noise.)

Exercise 1.11. Have a look at some of the properties that exhibited a Zipfian distribution in Bill, Juan, and Penousal's paper. See if you can implement them, and reproduce their results.